

An Introduction to High Performance Computing: Exercises

Paul Sumption
support@hpc.cam.ac.uk

Research Computing Services (<http://www.hpc.cam.ac.uk/>)
University Information Services (<http://www.uis.cam.ac.uk/>)

22nd May 2018 / CSD3 User Training

Exercise 1: Login

Using a Linux terminal you will login to the cluster with your HPC training account.

- ▶ Start the terminal by double clicking on the terminal icon
- ▶ In your terminal enter:
 - ▶ `ssh -Y abc123@login-cpu.hpc.cam.ac.uk`
Replace abc123 with your training account username
 - ▶ Enter your password as supplied on the sheet
 - ▶ Leave this terminal open, you will need it for exercise 3!

Exercise 2: Transfer some files

You will need to transfer the exercise files to the cluster.

- ▶ Open a second Linux terminal on your training computer.
- ▶ Enter this command: `cd ~\Course_material`
- ▶ Check the file 'exercises.tar' is in your directory listing
- ▶ Hint: `ls`

Exercise 3: Transfer some files

Transfer the exercises.tar to your HPC home folder.

- ▶ In the local terminal on your training computer enter the command:
 - ▶ `sftp abc123@login.hpc.cam.ac.uk`
Change abc123 to your training account username
 - ▶ The command: `put exercises.tar` will transfer the file from your local computer to the remote one
 - ▶ Type 'exit' to close the local terminal

Exercise 3: Learn more about a command

- (a) View the man page for the `cp` command by doing `man cp`. Use `SPACE` to page down and `b` to page up. Press `q` to exit the manual page command.
- (b) Copy `exercises.tar` to the `~/hpc-work` directory. Note that `~` is just a convenient shorthand for your home directory. Omitting the `~/` will look for a `hpc-work` in the current directory.
Hints: Do `cp exercises.tar ~/hpc-work/`. Note that you can often reduce typing by pressing `TAB`.

Exercise 4: Unzip the `exercises.tar` file

- (a) Use the `cd` command to enter the `~/hpc-work` directory and then list the contents — you should see the copy of `exercises.tgz`.
Hints: Do `cd ~/hpc-work/` then `ls -al`. Note that `cd ..` will take you back up one step to the home directory.
- (b) Unpack the tar archive to create an exercise subdirectory.
Hints: Do `tar -zxvf exercises.tar`

Exercise 5: File listings

- (a) In a terminal logged into the cluster list the contents of your current directory `ls`. This won't show everything — use `ls -al` for a long listing showing all files. Initially you will start in your home directory — use `pwd` to print the name of your current working directory. If you get lost, you can always do `cd` without arguments to return to your home directory.
- (b) Focus your long listing on **all files with names beginning “exercise”**.
Hints: Do `ls -al exercise*`
- (c) Print a long listing of the subdirectory `hpc-work`.
Hints: Do `ls -al hpc-work/`. Note that omitting the `/` reveals that the item `hpc-work` is actually a shortcut (technically a symbolic link) to `/hpc-work/username`.

Exercise 6: Environment Modules

- Connect to the cluster using your training account: See exercise 1 if you have closed your terminal.
- Get a list of modules that are currently loaded

Hints: `module list`

- Get a list of available R modules

Hints: `module av R`

Exercise 7: Run an Rscript

- ▶ Connect to the cluster using your training account: See exercise 1 if you have closed your terminal.
- ▶ In the exercises folder you transferred earlier there is a file called test.r
- ▶ Run this script using: Rscript hello.r
- ▶ Load the module for: R/3.3.2

Hints: [module load R/3.3.2](#)

- ▶ Run the script again: Rscript hello.r
- ▶ What happens? what changes?

Exercise 8: Explained

- ▶ Our R modules: module load r/(version). We have two sets of R modules, those with an upper case R where compiled for an older version of Linux and should be ignored (Darwin legacy)
- ▶ echo " " outputs the text between the quotes, >redirects the text into the .Renviro file.
- ▶ When we start R the .Renviro file is read and R will now be aware of our local library directory.
- ▶ .libPaths() is how to check your library locations

Exercise 8: Install the R library locally

As a user you can create a local R library directory for packages that you want to install.

- ▶ Load an R module: module load r-3.4.3-gcc-5.4.0-rbvnhga
- ▶ Create a folder in your home for your own R package installs: mkdir ~/my-R-libs
- ▶ Make R aware of the new library location:
echo "R_LIBS_USER=~ /my-R-libs" \textgreater;tee -a .Renviro
- ▶ Start R: R
- ▶ Display your library paths: .libPaths()
- ▶ Try loading a library: require(pander)
- ▶ Its not insalled, lets install it: install.packages("pander")
- ▶ Try loading a library: require(pander)
- ▶ Library is now installed, lets quit R: quit()

Exercise 9: Modules and Compilers

- ▶ Connect to the cluster using your training account: See exercise 1 if you have closed your terminal.
- ▶ Go to the [exercises](#) directory that you unzipped in hpcwork.
Try to compile the [hello.c](#) program using the default gcc compiler (it will fail because there is a deliberate bug).

Hints: [gcc hello.c -o hello](#)

- ▶ To fix the problem, open the [hello.c](#) file in the [gedit](#) editor.
Hints: Launch gedit in the background by doing [gedit&](#). A gedit window should appear. Remove the word [BUG](#), save the file and recompile. Do [./hello](#) to run the program.
- ▶ If you get this error:
WARNING **: cannot open display:
then you have missed the '-Y' in your SSH command

Exercise 9: Modules and Compilers

- ▶ The default version of gcc is 4.8.5. Compile hello.c again with `gcc 5.4.0`.

Hints: module av, module load gcc-5.4.0-gcc-4.8.5-fis24gg, then `gcc hello.c -o hello2`

Exercise 10: Submitting a Matlab job

- ▶ Submit a job which will run `matlab` on the `file.m` command file (which contains just the `ver` command).

Hints:

1. Load the matlab module using the `job_script` in your exercises directory.
2. Set the value of application to `"matlab -nodesktop -nosplash -nojvm"`
3. Set the value of options to `"-r file"`
4. Submit the job with `sbatch job_script`. The jobid is then printed.
5. Watch the job in the queue with `squeue`.
6. After it has disappeared, open the output file `slurm-jobid.out` in your editor. It should contain a list of licensed Matlab features.
7. For more demanding work you can increase the available memory by increasing the number of cpus.

Exercise 11: Submitting compiled code

- ▶ Submit a job which will run a copy of your hello program on 1 cpu.

Hints:

1. Edit the script `job_script` in your exercises directory. Set:
`#SBATCH --nodes=1`
`#SBATCH --ntasks=1`
`application="./hello"`
2. Submit the job with `sbatch job_script`. The jobid is then printed.
3. Watch the job in the queue with `squeue`.
4. After it has disappeared, open the output file `slurm-jobid.out` in your editor. There should be exactly one "Hello, World!" message.

Exercise 11: Array Jobs

- ▶ Submit your last job in the form of an array with indices 1-32. Use `-H` with `sbatch` to mark the array as held (so that it won't run immediately).

Hints:

1. Use `sbatch -H --array=1-32 job_script`
2. Use `squeue -u userid` to see your array job. Note that `-r` reports each array element individually.

- ▶ Release array element 1 and allow it to run. Then release the others.

Hints:

1. Use `scontrol release ${SLURM_ARRAY_JOB_ID}_1`
2. Use `squeue -u userid` again to watch what happens.
3. Release the others with
`scontrol release ${SLURM_ARRAY_JOB_ID}`
i.e. use the array id to release the entire array.
4. When all the jobs complete you should have 64 `slurm-${SLURM_ARRAY_JOB_ID}_N.out` files saying hello from various cpus on possibly multiple nodes.