

## An Introduction to High Performance Computing on the Minerva Cluster

Stuart Rankin  
sjr20@cam.ac.uk

Research Computing Services (<http://www.hpc.cam.ac.uk/>)  
University Information Services (<http://www.uis.cam.ac.uk/>)

24th April 2018 / NPL User Training

- ▶ Stuart Rankin — Research Computing User Services Team
- ▶ Paul Browne — Research Computing Platforms Team
- ▶ Course files can be downloaded from: [www.csd3.cam.ac.uk](http://www.csd3.cam.ac.uk)
- ▶ Please ask questions and let us know if you need assistance.

## Plan of the Course

Part 1: Basics

Part 2: HPC Facilities

Part 2: Using HPC

|             |         |
|-------------|---------|
| 10:00       | WELCOME |
| 11:30-11:45 | Break   |
| 13:00-14:00 | LUNCH   |
| 15:30-15:45 | Break   |
| 16:30       | CLOSE   |

## Basics: Why Buy a Big Computer?

What types of big problem might require a “Big Computer”?

*Compute Intensive:* A single problem requiring a large amount of computation.

*Memory Intensive:* A single problem requiring a large amount of memory.

*Data Intensive:* A single problem operating on a large amount of data.

*High Throughput:* Many unrelated problems to be executed in bulk.

## Basics: Scaling & Amdahl's Law

- ▶ Using more CPUs is not necessarily faster.
- ▶ Typically parallel codes have a *scaling limit*.
- ▶ Partly due to the system overhead of managing more copies, but also to more basic constraints;
- ▶ Amdahl's Law (idealized):

$$S(N) = \frac{1}{(1 - p + \frac{p}{N})}$$

where

$S(N)$  is the fraction by which the program has sped up relative to  $N = 1$


$p$  is the fraction of the program which can be parallelized

$N$  is the number of CPUs.

## Basics: Compute Intensive Problems

- ▶ Distribute the *work* for a *single problem* across multiple CPUs to reduce the execution time as far as possible.
- ▶ Program workload must be *parallelised*:
  - Parallel programs split into copies (processes or threads).
  - Each process/thread performs a part of the work on its own CPU, concurrently with the others.
  - A well-parallelised program will fully exercise as many CPUs as there are processes/threads.
- ▶ The CPUs typically need to exchange information rapidly, requiring specialized communication hardware.
- ▶ Many use cases from Physics, Chemistry, Engineering, Astronomy, Biology...
- ▶ The traditional domain of *HPC* and the *Supercomputer*.

## Basics: Amdahl's Law



<http://en.wikipedia.org/wiki/File:AmdahlsLaw.svg>

## The Bottom Line

- ▶ Parallelisation requires effort:
  - ▶ There are libraries to help (e.g. OpenMP, MPI).
  - ▶ First optimise performance on one CPU, then make  $p$  as large as possible.
- ▶ The scaling limit: eventually using more CPUs becomes detrimental instead of helpful.

## Basics: Data Intensive Problems

- ▶ Distribute the data for a single problem across multiple CPUs to reduce the overall execution time.
- ▶ The same work may be done on each data segment.
- ▶ Rapid movement of data to and from disk is more important than inter-CPU communication.
- ▶ Big Data problems of great current interest -
  - ▶ Hadoop/MapReduce
  - ▶ Life Sciences (genomics) and elsewhere.

## Basics: High Throughput

- ▶ Distribute independent, multiple problems across multiple CPUs to reduce the overall execution time.
- ▶ Workload is trivially (or embarrassingly) parallel:
  - \* Workload breaks up naturally into independent pieces.
  - \* Each piece is performed by a separate process/thread on a separate CPU (concurrently).
  - \* Little or no inter-CPU communication.
- ▶ Emphasis is on throughput over a period, rather than on performance on a single problem.
- ▶ Compute intensive capable  $\Rightarrow$  high throughput capable
- ▶ Compute intensive capable  $\nRightarrow$  high throughput capable

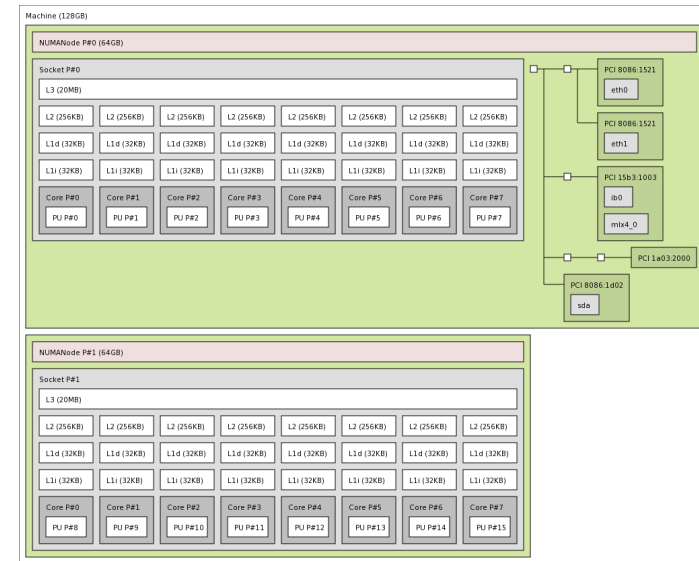
## Basics: Memory Intensive Problems

- ▶ Require aggregation of large memory into a single system image (i.e. a single computer running Linux).
- ▶ Technically more challenging to build machines (very fast, low latency interconnection between all CPUs and all memory).
- ▶ Coding/porting easier (memory appears seamless).
- ▶ Optimisation harder (memory is actually highly nonuniform).
- ▶ Historically, the arena of large SGI systems.
- ▶ Nowadays, similar complexity inside single commodity servers.

## Basics: Inside a Modern Computer

- ▶ Today's commodity servers already aggregate both CPUs and memory to make a single system image in a single box.
- ▶ Even small computers now have multiple CPU cores per socket  
 $\Rightarrow$  each socket is a Symmetric Multi-Processor (SMP).
- ▶ Larger computers have multiple sockets (each with local memory):  
 all CPU cores (unequally) share the node memory  
 $\Rightarrow$  the node is a shared memory multiprocessor with Non-Uniform Memory Architecture (NUMA)  
 but users still see a single computer (single system image).

## Basics: Inside a Modern Computer




## Basics: How to Build a Supercomputer

- ▶ A supercomputer aggregates contemporary CPUs and memory to obtain increased computing power.
  - ▶ Usually today these are clusters.
1. Take some (multicore) CPUs plus some memory.
    - ▶ Could be an off-the-shelf server, or something more special.
    - ▶ A NUMA, shared memory, multiprocessor building block: a node.

## Basics: How to Build a Supercomputer

2. Connect the nodes with one or more networks. E.g.  
 Gbit Ethernet: 100 MB/sec  
 Omni-Path: 10 GB/sec

Faster network is for inter-CPU communication across nodes.  
 Slower network is for management and provisioning.  
 Storage may use either.



### 3. Logically bind the nodes

- ▶ Clusters consist of distinct nodes (i.e. separate Linux computers) on common private network(s) and controlled centrally.
  - \* Private networks allow CPUs in different nodes to communicate.
  - \* Clusters are **distributed memory** machines:  
Each process/thread sees only its local node's CPUs and memory (without help).
  - \* **Each process/thread must fit within a single node's memory.**
- ▶ More expensive machines logically bind nodes into a single system i.e. CPUs and memory.
  - \* E.g. SGI UV.
  - \* Private networks allow CPUs to see CPUs and memory in other nodes.
  - \* These are **shared memory** machines.
  - \* Logically a single system - 1 big node - but very non-uniform.
  - \* A single process can span the entire system.

- ▶ Non-parallel (serial) code
  - \* For a single node as for a workstation.
  - \* Typically run as many copies per node as CPUs, assuming node memory is sufficient.
  - \* Replicate across multiple nodes.
- ▶ Parallel code
  - \* Shared memory methods within a node.  
E.g. pthreads, OpenMP.
  - \* Distributed memory methods spanning multiple nodes.  
Message Passing Interface (MPI).

## Basics: Summary

- ▶ Why have a supercomputer?
  - ▶ Big single problems, many problems, Big Data.
- ▶ Most current supercomputers are **clusters** of separate **nodes**.
- ▶ Each node has **multiple CPUs** and **non-uniform shared memory**.
- ▶ **Parallel code** uses shared memory (pthreads/OpenMP) within a node, distributed memory (MPI) spanning multiple nodes.
- ▶ **Non-parallel code** uses the memory of one node, but may be copied across many.

- ▶ Each compute node:
  - \* 2x16 cores, Intel Skylake 2.6 GHz **32 CPUs**
  - \* 192 GB RAM **6 GB per CPU**
  - \* 100 Gb/sec Omni-Path **10 GB/sec (for MPI and storage)**
- ▶ 16 compute nodes
- ▶ 1 high-memory login/head node ([minerva-login1.npl.co.uk](http://minerva-login1.npl.co.uk)), 768 GB RAM.
- ▶ 1 GPU node, 2 × NVIDIA P100 GPU, 192 GB RAM.

- ▶ Cambridge Service for Data Driven Discovery
- ▶ Peta4 — Intel CPU cluster
- ▶ Wilkes2 — NVIDIA GPU cluster
- ▶ Hadoop-based data analytic platform
- ▶ Burst buffer
- ▶ Some NPL users through CORE.

- ▶ Each compute node:
  - \* 2x16 cores, Intel Skylake 2.6 GHz **32 CPUs**
  - \* 192 GB or 384 GB RAM **6 GB or 12 GB per CPU**
  - \* 100 Gb/sec Omni-Path **10 GB/sec (for MPI and storage)**
- ▶ 768 compute nodes
- ▶ 8 login nodes ([login-cpu.hpc.cam.ac.uk](http://login-cpu.hpc.cam.ac.uk))

- ▶ CPUs are **general purpose**
- ▶ Some types of parallel workload fit **vector** processing well:
  - ▶ Single Instruction, Multiple Data (SIMD)
  - ▶ *Think pixels on a screen*
  - ▶ GPUs specialise in this type of work
  - ▶ Also competitor many-core architectures such as the Intel Phi

## Wilkes2-GPU

- ▶ Each compute node:
  - \* 4 × NVIDIA P100 GPU4 GPUs
  - \* 1x12 cores, Intel Broadwell 2.2 GHz12 CPUs
  - \* 96 GB RAM96 GB RAM
  - \* 100 Gb/sec (4X EDR) Infiniband.10 GB/sec (for MPI and storage)
- ▶ 90 compute nodes.
- ▶ 8 login nodes ([login-gpu.hpc.cam.ac.uk](http://login-gpu.hpc.cam.ac.uk)).

## Peta4-KNL (Intel Phi)

- ▶ Each compute node:
  - \* 64 cores, Intel Phi 7210256 CPUs
  - \* 96 GB RAM96 GB RAM
  - \* 100 Gb/sec Omni-Path10 GB/sec (for MPI and storage)
- ▶ 342 compute nodes
- ▶ Shared login nodes with Peta4-Skylake

## Cluster Storage

- ▶ Minerva uses NFS to share user directories between all cluster nodes (150 TB).
- ▶ CSD3 uses the Lustre cluster filesystem:
  - \* Very scalable, high bandwidth.
  - \* Multiple RAID6 back-end disk volumes.
  - \* Multiple object storage servers.
  - \* Single metadata server.
  - \* Tape-backed HSM on newest filesystems.
  - \* 12 GB/sec overall read or write.
  - \* Prefers big read/writes over small.

## Obtaining an Account and Support

- ▶ Please contact the NPL IT Service Desk:
  - ▶ [itservicedesk@npl.co.uk](mailto:itservicedesk@npl.co.uk)
  - ▶ Room F12-CS1
  - ▶ Tel. 6000
- ▶ Second line support is provided by the Cambridge RCS.

## Part III: Using HPC

### Connecting: Windows Clients

- ▶ putty, pscp, psftp  
<http://www.chiark.greenend.org.uk/~sgtatham/putty/download.html>
- ▶ WinSCP  
<http://winscp.net/eng/download.php>
- ▶ TurboVNC (remote desktop, 3D optional)  
<http://sourceforge.net/projects/turbovnc/files/>
- ▶ Cygwin (provides an application environment similar to Linux)  
<http://cygwin.com/install.html>  
Includes X server for displaying graphical applications remotely.
- ▶ MobaXterm  
<http://mobaxterm.mobatek.net/>

### Using HPC: Connecting

- ▶ SSH secure protocol only.  
Supports login, file transfer, remote desktop. . .

### Connecting: Linux/MacOSX/UNIX Clients

- ▶ **ssh**, scp, sftp, **rsync**  
Installed (or installable).
- ▶ TurboVNC (remote desktop, 3D optional)  
<http://sourceforge.net/projects/turbovnc/files/>
- ▶ On MacOSX, install XQuartz to display remote graphical applications.  
<http://xquartz.macosforge.org/landing/>



## Connecting: Login

- ▶ From graphical clients:  
Host: `minerva-login1.npl.co.uk`  
Username: `npl\abc123` (your NPL AD account name)
- ▶ From Linux/MacOSX/UNIX (or Cygwin):  
`ssh -Y npl\\abc12@minerva-login1.npl.co.uk`  
  
Note the double backslash — this is because UNIX command interpreters treat `\` as special.

## Connecting: First time login

- ▶ The first connection to a particular hostname produces the following:  

```
The authenticity of host 'minerva-login1.npl.co.uk (139.143.201.10)' can't be established.  
  
ECDSA key fingerprint is SHA256:k/eB+LjcAfQW56XCzK9QptT0wVWF7j3a/CPxPRd7+1E.  
ECDSA key fingerprint is MD5:18:9a:97:e2:87:4c:07:60:cb:43:46:f2:bb:d8:3d:01.  
  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added 'minerva-login1.npl.co.uk (139.143.201.10)' (ECDSA) to the list of known hosts.
```
- ▶ One should always check the fingerprint before typing “yes”.
- ▶ Graphical SSH clients *should* ask a similar question.
- ▶ Designed to detect fraudulent servers.

## Connecting: First time login

- ▶ Exercise 1 - Log into your Minerva account.
- ▶ Exercise 2 - Simple command line operations.

## Connecting: File Transfer

- ▶ With graphical clients, connect as before and drag and drop.
- ▶ From Linux/MacOSX/UNIX (or Cygwin):  

```
rsync -av old_directory/  
npl\\abc12@minerva-login1.npl.co.uk:hpc-work/new_directory  
copies contents of old_directory to ~/hpc-work/new_directory.  
  
rsync -av old_directory  
npl\\abc12@minerva-login1.npl.co.uk:hpc-work/new_directory  
copies old_directory (and contents) to  
~/hpc-work/new_directory/old_directory.  
  
* Rerun to update or resume after interruption.  
* All transfers are checksummed.  
* For transfers in the opposite direction, place the remote machine as  
the first argument.
```
- ▶ Exercise 3 - File transfer.

## Connecting: Remote Desktop

- First time starting a remote desktop:

```
[sjr20@login-a-1 ~]$ vncserver
```

You will require a password to access your desktops.

Password:

Verify:

Would you like to enter a view-only password (y/n)? n

New 'login-a-1:99 (sjr20)' desktop is login-a-1:99

Starting applications specified in /home/sjr20/.vnc/xstartup  
Log file is /home/sjr20/.vnc/login-a-1:99.log

- NB Choose a **different** password for VNC.
- The VNC password protects your desktop from other users.
- Remember the unique display number (**99** here) of your desktop.

## Connecting: Remote Desktop

- Remote desktop already running:

```
[sjr20@login-a-1 ~]$ vncserver -list
```

TigerVNC server sessions:

| X DISPLAY # | PROCESS ID |
|-------------|------------|
| :99         | 130655     |

- Kill it:

```
[sjr20@login-a-1 ~]$ vncserver -kill :99  
Killing Xvnc process ID 130655
```

- Typically you only need **one** remote desktop.
- Keeps running until killed, or the node reboots.

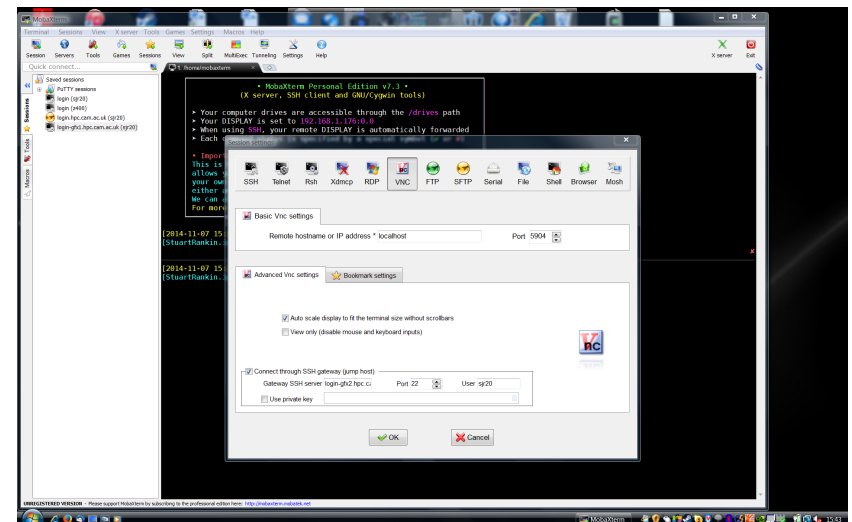
## Connecting: Remote Desktop

- To connect to the desktop from Linux:

```
vncviewer -via npl\\abc12@minerva-login1.npl.ad.local localhost:99
```

- The display number **99** will be different in general and unique to each desktop.
- You will be asked firstly for your AD login password, and secondly for your VNC password.
- Press **F8** to bring up the control panel.
- Exercise 4 - Remote desktop (from Windows)

## Connecting: Remote Desktop (MobaXterm)



- ▶ CentOS Linux 7.4 (Red Hat Enterprise Linux 7.4 rebuild)
  - ▶ bash shell
  - ▶ Gnome or XFCE4 desktop (if you want)
  - ▶ GCC compilers and other development software.
- ▶ But you don't need to know that.

- ▶ `/home/abc12@npl.ad.local`
  - ▶ 50GB quota.
  - ▶ Visible equally from all nodes.
  - ▶ Single storage server.
  - ▶ Regular backups.
  - ▶ Not intended for job outputs or large/many input files.
- ▶ `~/hpc-work`
  - ▶ Visible equally from all nodes.
  - ▶ Larger (1TB initial quota).
  - ▶ Intended for job inputs and outputs.
  - ▶ **Not backed up by default.**

- ▶ `quota`

```
[sjr20@login-a-1 ~]$ quota -s
Disk quotas for user sjr20 (uid 1004):
Filesystem space quota limit grace files quota limit grace
10.44.82.252:/hpc-work
OK 1024G 1126G 1 0 0
10.44.82.252:/home
13272K 51200M 56320M 345 0 0
```

- ▶ Aim to stay below the soft limit (*quota*).
- ▶ Once over the soft limit, you have 7 days grace to return below.
- ▶ When the grace period expires, or you reach the hard limit (*limit*), no more data can be written.
- ▶ It is important to rectify an out of quota condition ASAP.

- ▶ `quota`

```
[sjr20@login-a-1 ~]$ quota -s
Disk quotas for user sjr20 (uid 1004):
Filesystem space quota limit grace files quota limit grace
10.44.82.252:/hpc-work
OK 1024G 1126G 1 0 0
10.44.82.252:/home
13272K 51200M 56320M 345 0 0
```

- ▶ Aim to stay below the soft limit (*quota*).
- ▶ Once over the soft limit, you have 7 days grace to return below.
- ▶ When the grace period expires, or you reach the hard limit (*limit*), no more data can be written.
- ▶ It is important to rectify an out of quota condition ASAP.

- ▶ **Be careful and if unsure, please ask support.**
  - ▶ Can lead to **accidental destruction** of your data or **account compromise**.
- ▶ Avoid changing the permissions on your home directory.
  - ▶ Files under `/home` are particularly security sensitive.
  - ▶ Easy to break passwordless communication between nodes.

- ▶ Free software accompanying **Red Hat Enterprise Linux** is (or can be) provided.
- ▶ Other software (free and non-free) is available via **modules**.
- ▶ Proprietary software currently available includes Matlab and COMSOL.
- ▶ New software may be possible to provide on request.
- ▶ **Self-installed software should be properly licensed.**
- ▶ ***sudo will not work.*** (You should be worried if it did.)

- ▶ Modules load or unload additional software packages.
- ▶ Some are **required** and automatically loaded on login.
- ▶ Others are optional extras, or possible replacements for other modules.
- ▶ **Beware** unloading default modules in `~/.bashrc`.
- ▶ **Beware** overwriting environment variables such as `PATH` and `LD_LIBRARY_PATH` in `~/.bashrc`. If necessary append or prepend.

- ▶ Currently loaded:

```
module list
Currently Loaded Modulefiles:
  1) dot                      3) centos7/global
  2) slurm                    4) centos7/default-basic
```

- ▶ Available:

```
module av
```

## User Environment: Environment Modules

### ► Whatis:

```
module whatis openmpi-3.0.0-gcc-4.8.5-n2hvjgm
openmpi-3.0.0-gcc-4.8.5-n2hvjgm: The Open MPI Project is an open source...
```

### ► Load:

```
module load openmpi-3.0.0-gcc-4.8.5-n2hvjgm
```

### ► Unload:

```
module unload openmpi-3.0.0-gcc-4.8.5-n2hvjgm
```

## User Environment: Environment Modules

### ► Matlab

```
module load matlab/r2018a
```

### ► Invoking matlab in batch mode:

```
matlab -nodisplay -nojvm -nosplash command
```

where the file `command.m` contains your matlab code.

### ► The current site license contains the Parallel Computing Toolbox.

## User Environment: Environment Modules

### ► Purge:

```
module purge
```

### ► Defaults:

```
module show centos7/default-basic
module load centos7/default-basic
```

### ► Run time environment must match compile time environment.

## User Environment: Compilers

### ► GCC

```
gcc -O3 -mtune=native code.c -o prog
gfortran -O3 -mtune=native code.f90 -o prog
```

```
module load openmpi-3.0.0-gcc-4.8.5-n2hvjgm
mpicc -O3 -mtune=native mpi_code.c -o mpi_prog
mpif90 -O3 -mtune=native mpi_code.f90 -o mpi_prog
```

### ► Exercise 5: Modules and Compilers



- ▶ Compute resources are managed by a scheduler:  
[SLURM](#)/[PBS](#)/[SGE](#)/[LSF](#)/...
- ▶ Jobs are submitted to the scheduler
  - analogous to submitting jobs to a print queue
  - a file (*submission script*) is copied and queued for processing.

- ▶ Jobs are submitted from the [login node](#)
  - not itself managed by the scheduler.
- ▶ Jobs may be either non-interactive ([batch](#)) or [interactive](#).
- ▶ [Batch](#) jobs run a shell script on the first of a list of allocated nodes.
- ▶ [Interactive](#) jobs provide a command line on the first of a list of allocated nodes.

- ▶ Jobs may use [part](#) or [all](#) of one or more nodes
  - the owner can specify `--exclusive` to force exclusive node access.
- ▶ Template submission scripts are available under [~/job\\_templates](#).

## Job Submission: Using SLURM

- Prepare a shell script and submit it to SLURM:

```
[abc123@login-a-1]$ sbatch slurm_submission_script
Submitted batch job 790299
```

## Job Submission: Show Queue

- Submitted job scripts are copied and stored in a queue:

```
[abc123@login-a-1]$ squeue -u abc123
      JOBID PARTITION   NAME     USER ST       TIME  NODES NODELIST(REASON)
      790299   skylake    Test3    abc123 PD        0:00        2 (PriorityResourcesAssocGrpCPUMinsLimit)
      790290   skylake    Test2    abc123 R       27:56:10        2 cpu-a-[1,10]
```

## Job Submission: Monitor Job

- Examine a particular job:

```
[abc123@login-a-1]$ scontrol show job=790290
```

## Job Submission: Cancel Job

- Cancel a particular job:

```
[abc123@login-a-1]$ scancel 790290
```

## Job Submission: Scripts

### ► SLURM

In `~/job_templates`, see examples: `slurm_submit.skylake.generic`, `slurm_submit.skylake.matlab`.

```
#!/bin/bash
#! Name of the job:
#SBATCH -J myjob
#! Which project should be charged:
#SBATCH -A NPL-GENERAL-CPU
#! How many whole nodes should be allocated?
#SBATCH --nodes=1
#! How many tasks will there be in total? (<= nodes*32)
#SBATCH --ntasks=116
#! How much wallclock time will be required?
#SBATCH --time=02:00:00
#! Select partition:
#SBATCH -p skylake
...
```

- **#SBATCH** lines are *structured comments*
  - correspond to sbatch command line options.
- The above job will be given **1 cpu16 cpus** on 1 node for 2 hours (by default there is 1 task per node, and 1 cpu per task).

## Job Submission: Accounting Commands

### ► How many core hours available do I have?

mybalance

| User  | Usage | Account     | Usage | Account Limit | Available (hours) |
|-------|-------|-------------|-------|---------------|-------------------|
| ----- | ----- | -----       | ----- | -----         | -----             |
| sjr20 | 3     | SUPPORT-CPU | 2,929 | 22,425,600    | 22,422,671        |
| sjr20 | 0     | SUPPORT-GPU | 0     | 87,600        | 87,600            |

### ► How many core hours does some other project or user have?

gbalance -p SUPPORT-CPU

| User               | Usage | Account     | Usage | Account Limit | Available (hours) |
|--------------------|-------|-------------|-------|---------------|-------------------|
| -----              | ----- | -----       | ----- | -----         | -----             |
| pfb29              | 2,925 | SUPPORT-CPU | 2,929 | 22,425,600    | 22,422,671        |
| sjr20 *            | 3     | SUPPORT-CPU | 2,929 | 22,425,600    | 22,422,671        |
| ...                |       |             |       |               |                   |
| (Use -u for user.) |       |             |       |               |                   |

### ► List all jobs charged to a project/user between certain times:

| gstatment -p NPL-GENERAL-CPU -u xyz10 -s "2018-04-01-00:00:00" -e "2018-04-30-00:00:00" |       |            |            |           |                     |          |            |         |  |
|-----------------------------------------------------------------------------------------|-------|------------|------------|-----------|---------------------|----------|------------|---------|--|
| JobID                                                                                   | User  | Account    | JobName    | Partition | End                 | ExitCode | State      | CompHrs |  |
| 263                                                                                     | xyz10 | support-c+ | _interact+ | skylake   | 2018-04-18T19:44:40 | 0:0      | TIMEOUT    | 1.0     |  |
| 264                                                                                     | xyz10 | support-c+ | _interact+ | skylake   | 2018-04-18T19:48:07 | 0:0      | CANCELLED+ | 0.1     |  |
| 275                                                                                     | xyz10 | support-c+ | _interact+ | skylake   | Unknown             | 0:0      | RUNNING    | 0.3     |  |
| 475                                                                                     |       |            |            |           |                     |          |            |         |  |

## Job Submission: Single Node Jobs

### ► Serial jobs requiring large memory, or OpenMP codes.

```
#!/bin/bash
...
#SBATCH --nodes=1
#SBATCH --ntasks=1
# Default is 1 task per node
#SBATCH --cpus-per-task=

#SBATCH --mem=5990
# Memory per node in MB - default is pro rata by cpu number
# Increasing --mem or --cpus-per-task implicitly increases the other
...
export OMP_NUM_THREADS= # For OpenMP across cores
$application $options
...
```

## Job Submission: Single Node Jobs

### ► Serial jobs requiring large memory, or OpenMP codes.

```
#!/bin/bash
...
#SBATCH --nodes=1
#SBATCH --ntasks=1
# Default is 1 task per node
#SBATCH --cpus-per-task=1
# Default is 1 cpu (core) per task
#SBATCH --mem=5990
# Memory per node in MB - default is pro rata by cpu number
# Increasing --mem or --cpus-per-task implicitly increases the other
...
export OMP_NUM_THREADS= # For OpenMP across cores
$application $options
...
```



## Job Submission: Single Node Jobs

- ▶ Serial jobs requiring large memory, or OpenMP codes.

```
#!/bin/bash
...
#SBATCH --nodes=1
#SBATCH --ntasks=1
# Default is 1 task per node
#SBATCH --cpus-per-task=32 # Whole node

#SBATCH --mem=5990
# Memory per node in MB - default is pro rata by cpu number
# Increasing --mem or --cpus-per-task implicitly increases the other
...
export OMP_NUM_THREADS=32 # For OpenMP across 32 cores
$application $options
...
```

## Job Submission: Single Node Jobs

- ▶ Serial jobs requiring large memory, or OpenMP codes.

```
#!/bin/bash
...
#SBATCH --nodes=1
#SBATCH --ntasks=1
# Default is 1 task per node
#SBATCH --cpus-per-task=16 # Half node

#SBATCH --mem=5990
# Memory per node in MB - default is pro rata by cpu number
# Increasing --mem or --cpus-per-task implicitly increases the other
...
export OMP_NUM_THREADS=16 # For OpenMP across 16 cores
$application $options
...
```

## Job Submission: Single Node Jobs

- ▶ Serial jobs requiring large memory, or OpenMP codes.

```
#!/bin/bash
...
#SBATCH --nodes=1
#SBATCH --ntasks=1
# Default is 1 task per node
#SBATCH --cpus-per-task=32 # Whole node

#SBATCH --mem=5990
# Memory per node in MB - default is pro rata by cpu number
# Increasing --mem or --cpus-per-task implicitly increases the other
...
export OMP_NUM_THREADS=16 # For OpenMP across 16 cores (using all memory)
$application $options
...
```

## Job Submission: MPI Jobs

- ▶ Parallel job across multiple nodes.

```
#!/bin/bash
...
#SBATCH --nodes=4
#SBATCH --ntasks=128 # i.e. 32x4 MPI tasks in total.
#SBATCH --cpus-per-task=2
...
mpirun -np 128 $application $options
...
```

- ▶ SLURM-aware MPI launches remote tasks via SLURM.
- ▶ The template script uses \$SLURM\_TASKS\_PER\_NODE to set PPN.

## Job Submission: MPI Jobs

- ▶ Parallel job across multiple nodes.

```
#!/bin/bash
...
#SBATCH --nodes=4
#SBATCH --ntasks=64      # i.e. 16x4 MPI tasks in total.
#SBATCH --cpus-per-task=2
...
mpirun -ppn 16 -np 64 $application $options
...
```

- ▶ SLURM-aware MPI launches remote tasks via SLURM.
- ▶ The template script uses \$SLURM\_TASKS\_PER\_NODE to set PPN.

## Job Submission: Hybrid Jobs

- ▶ Parallel jobs using both MPI and OpenMP.

```
#!/bin/bash
...
#SBATCH --nodes=4
#SBATCH --ntasks=64      # i.e. 16x4 MPI tasks in total.
#SBATCH --cpus-per-task=2
...
export OMP_NUM_THREADS=2  # i.e. 2 threads per MPI task.
mpirun -ppn 16 -np 64 $application $options
...
```

- ▶ This job uses 128 CPUs (each MPI task splits into 2 OpenMP threads).

## Job Submission: High Throughput Jobs

- ▶ Multiple serial jobs across multiple nodes.
- ▶ Use `srun` to launch tasks (job steps) within a job.

```
#!/bin/bash
...
#SBATCH --nodes=2
...
cd directory_for_job1
srun --exclusive -N 1 -n 1 $application $options_for_job1 > output 2> err &
cd directory_for_job2
srun --exclusive -N 1 -n 1 $application $options_for_job2 > output 2> err &
...
cd directory_for_job64
srun --exclusive -N 1 -n 1 $application $options_for_job64 > output 2> err &
wait
```

- ▶ Exercise 6 - Submitting Jobs.

## Job Submission: Interactive

- ▶ Compute nodes are accessible via SSH while you have a job running on them.
- ▶ Alternatively, submit an interactive job:  
`sintr -A NPL-GENERAL-CPU -N1 -n8 -t 2:0:0`
- ▶ Within the window (screen session):
  - \* Launches a shell on the first node (when the job starts).
  - \* Graphical applications should display correctly.
  - \* Create new shells with `ctrl-a c`, navigate with `ctrl-a n` and `ctrl-a p`.
  - \* `ssh` or `srun` can be used to start processes on any nodes in the job.
  - \* SLURM-aware MPI will do this automatically.

## Job Submission: Array Jobs

- ▶ [http://slurm.schedmd.com/job\\_array.html](http://slurm.schedmd.com/job_array.html)
- ▶ Used for submitting and managing large sets of similar jobs.
- ▶ Each job in the array has the same initial options.
- ▶ SLURM

```
[abc123@login-a-1]$ sbatch --array=1-7:21,3,5,7 -A NPL-GENERAL-CPU submit_script
Submitted batch job 791609
```

```
[abc123@login-a-1]$ squeue -u abc123
```

| JOBID    | PARTITION | NAME | USER   | ST | TIME | NODES | NODELIST(REASON) |
|----------|-----------|------|--------|----|------|-------|------------------|
| 791609_1 | skylake   | hpl  | abc123 | R  | 0:06 | 1     | cpu-a-6          |
| 791609_3 | skylake   | hpl  | abc123 | R  | 0:06 | 1     | cpu-a-16         |
| 791609_5 | skylake   | hpl  | abc123 | R  | 0:06 | 1     | cpu-a-7          |
| 791609_7 | skylake   | hpl  | abc123 | R  | 0:06 | 1     | cpu-a-7          |

791609\_1, 791609\_3, 791609\_5, 791609\_7

i.e.  $\{\text{SLURM\_ARRAY\_JOB\_ID}\}_{\text{SLURM\_ARRAY\_TASK\_ID}}$

SLURM\_ARRAY\_JOB.ID = SLURM\_JOBID for the first element.

## Job Submission: Array Jobs (ctd)

- ▶ Updates can be applied to specific array elements using  $\{\text{SLURM\_ARRAY\_JOB\_ID}\}_{\text{SLURM\_ARRAY\_TASK\_ID}}$
- ▶ Alternatively operate on the entire array via  $\{\text{SLURM\_ARRAY\_JOB\_ID}\}$ .
- ▶ Some commands still require the SLURM\_JOB.ID (sacct, sreport, sshare, sstat and a few others).
- ▶ Exercise 7 - Array Jobs.

## Scheduling

- ▶ SLURM scheduling is multifactor:
  - ▶ QoS — payer or non-payer?
  - ▶ Age — how long has the job waited?  
Don't cancel jobs that seem to wait too long.
  - ▶ Fair Share — how much recent usage?  
Payers with little recent usage receive boost (not implemented yet).
  - ▶ sprio -j jobid
- ▶ Backfilling
  - ▶ Promote lower priority jobs into gaps left by higher priority jobs.
  - ▶ Demands that the higher priority jobs not be delayed.
  - ▶ Relies on reasonably accurate wall time requests for this to work.
  - ▶ Jobs of default length will not backfill readily.

## Wait Times

- ▶ 36 hour job walltimes are permitted.
- ▶ This sets the timescale at busy times (*without* backfilling).
- ▶ Use backfilling when possible.
- ▶ Short (1 hour or less) jobs have higher throughput.

- ▶ Insurance against failures during long jobs.
- ▶ Restart from checkpoints to work around finite job length.
- ▶ Application native methods are best. Failing that ...

- ▶ **Do ...**
  - ▶ Give reasonably accurate wall times (allows [backfilling](#)).
  - ▶ Check your balance occasionally ([mybalance](#)).
  - ▶ Test on a small scale first.
  - ▶ Implement [checkpointing](#) if possible (reduces resource wastage).
- ▶ **Don't ...**
  - ▶ Request more than you need
    - you will wait longer and use more credits.
  - ▶ Cancel jobs unnecessarily
    - priority increases over time.